


**for\_each** in tf  
iterate over maps or sets  
if you need to dyanmically  
create IAM users or s3  
buckets

```
variable "users" {  
  type = map(string)  
  default = {  
    alice = "Admin"  
    bob   = "Developer"  
    charlie = "Tester"  
  }  
}  
  
resource "aws_iam_user" "users" {  
  for_each = var.users  
  name     = each.key  
  tags = {  
    Role = each.value  
  }  
}
```



**Scenario:**  
You are a **DevOps Engineer** at a growing startup. Your team wants to deploy an **infrastructure-as-code (IaC) solution** for setting up the following AWS resources **across multiple environments (dev, staging, prod)** using **Terraform Workspaces, Modules, Variables, and Outputs**.

- Requirements:**
- 1 **Use Terraform Workspaces** to differentiate between **dev, staging, and prod** environments.
  - 2 **Create a reusable Terraform module** for EC2 instances and S3 buckets.
  - 3 The **EC2 instance** should:
    - Use Amazon Linux 2
    - Have an SSH security group
    - Attach an IAM role to allow S3 access
    - 4 The **S3 bucket** should:
      - Have versioning enabled
      - Be environment-specific (my-app-dev, my-app-staging, etc.)
    - 5 **Use input variables** for instance type, region, and S3 bucket name.
    - 6 **Use Terraform outputs** to display EC2 public IP and S3 bucket name.

- Task:**
- ✓ **Step 1:** Create a Terraform module (ec2\_s3\_module) that:
    - Creates an **EC2 instance** with an IAM role allowing S3 access.
    - Creates an **S3 bucket** with versioning enabled.
  - ✓ **Step 2:** Use **Terraform Workspaces** to deploy this module for **dev, staging, and prod** environments.
  - ✓ **Step 3:** Use **input variables** for EC2 instance type, region, and S3 bucket name.
  - ✓ **Step 4:** Use **Terraform Outputs** to show:
    - EC2 **public IP**
    - S3 **bucket name**